

Inferring Occupancy Duration from Indoor CO₂ Concentrations in University Dormitory Rooms: A Block-Level Excess Approach with Monte Carlo Uncertainty Quantification

Gautam Author^{a,*}

^aCenter for Research in Indoor and Personal Technologies (CRIPT), , , , USA

Abstract

Estimating how long occupants spend in their rooms is valuable for understanding indoor environmental exposure, energy use, and well-being in residential buildings. While CO₂ concentration is widely used as a proxy for occupancy detection (present/absent), estimating *duration* of occupancy within time blocks remains an open challenge. We present a physics-based method that infers occupied minutes within four-hour blocks from indoor CO₂ measurements in university dormitory rooms. The method fits an autoregressive AR(1) model to each sensor's CO₂ time series, calibrates block-level excess CO₂ against per-sensor empty-room behavior, and converts the calibrated signal to occupied minutes using finite-window mass balance physics. Monte Carlo propagation of parameter uncertainty yields 80% credible intervals for each estimate. We validate the method on 12 sensors monitoring 24 subjects across 774 four-hour blocks over 68 days. All four independent evaluation criteria—covering cross-sensor generalizability (E1–E2), parameter sensitivity (E3), and uncertainty informativeness (E4)—are satisfied. Two calibration diagnostics used to tune the empty-room floor are met by construction. An out-of-sample temporal split shows that the calibration partially generalizes but degrades when training data are limited. A semisynthetic validation—injecting known occupancy patterns into real empty-room CO₂ data with model-mismatch stress tests—provides partial external validation of duration recovery accuracy. Comparison against threshold and slope baseline methods demonstrates that the block-excess approach achieves superior empty-room specificity. The method requires no labeled minute-level training data, uses only CO₂ measurements and known-empty calibration periods, and runs in under one minute on standard hardware.

Keywords: occupancy duration, CO₂ inference, indoor air quality, autoregressive model, Monte Carlo uncertainty, dormitory buildings

*Corresponding author

Email address: gautam@university.edu (Gautam Author)

1. Introduction

Understanding how long building occupants spend in particular spaces is fundamental to indoor environmental quality (IEQ) assessment, energy management, and occupant well-being research [1]. In residential buildings—particularly university dormitories where students live, study, and sleep—the duration of room occupancy directly determines cumulative exposure to indoor pollutants, thermal conditions, and lighting environments [2, 3]. Occupancy-driven energy models further rely on accurate time-use data to predict heating, cooling, and ventilation loads [4, 5].

While numerous methods exist for binary occupancy detection (present vs. absent), far fewer address the more informative question of *how long* an occupant was present during a given time window. Binary detection from CO₂ thresholds is well established [6, 7, 8], and multi-sensor fusion approaches combining CO₂ with temperature, humidity, and light have pushed detection accuracy above 95% [9]. Machine learning approaches—including random forests, support vector machines, and deep recurrent networks—have improved detection accuracy further [10, 11, 12] but typically still target binary or count outcomes rather than continuous duration. Bayesian and state-space formulations have been applied to CO₂-based occupancy counting [13, 14], yet these also estimate the number of people present rather than how long they stay.

CO₂ concentration is a particularly attractive proxy for occupancy in dormitory settings because it is noninvasive, inexpensive to measure continuously, and directly linked to human metabolic activity through respiration. The physics of indoor CO₂ are well understood: occupants generate CO₂ at rates proportional to their metabolic activity (12 L/h to 20 L/h per person at rest [15]), while ventilation removes CO₂ at a rate proportional to the concentration difference between indoor and outdoor air [16]. This mass-balance relationship creates a direct, physics-based link between CO₂ dynamics and occupancy patterns and underpins demand-controlled ventilation strategies [17].

However, extracting *duration* information from CO₂ time series faces several challenges. First, the decay time constant of indoor CO₂ (typically 30–120 minutes in naturally ventilated rooms) means that CO₂ remains elevated long after occupants depart, creating ambiguity between current and residual occupancy signal. Second, sensor noise, outdoor CO₂ fluctuations, and building-level ventilation transients introduce confounding variation. Third, the generation rate per occupant depends on metabolic activity, which varies with activity type and individual physiology.

In this paper, we present a method that addresses these challenges by operating at the *block level* (four-hour windows) rather than attempting minute-by-minute occupancy reconstruction. This design choice is motivated by a fundamental signal-to-noise ratio (SNR) argument: at the minute timescale, the occupancy-driven CO₂ innovation (typically 5 ppm/min to 60 ppm/min) is comparable to or smaller than sensor noise (5 ppm to 40 ppm), yielding minute-level SNR of order unity. At the block level, averaging over 240 minutes reduces *innovation* noise, though the AR(1) autocorrelation ($\phi \approx 0.995$) limits the effec-

tive noise reduction compared to the i.i.d. case; the practical gain comes from aggregating the cumulative CO₂ signal over a long window rather than from simple \sqrt{n} averaging.

The contributions of this paper are:

1. A physics-based method for estimating occupancy *duration* (not just presence/absence) from CO₂, requiring no labeled minute-level training data—only binary block-level presence labels and known-empty calibration periods.
2. Per-sensor adaptive calibration that accounts for room-specific ventilation, sensor placement, and residual CO₂ characteristics.
3. Full Monte Carlo uncertainty quantification providing 80% credible intervals for each block-level duration estimate.
4. A structured validation framework distinguishing calibration diagnostics from independent evaluation criteria for CO₂-based duration estimation.
5. Application to 24 university dormitory residents across 12 rooms over 68 days, demonstrating feasibility in real-world conditions.

Fig. 1 provides an overview of the proposed method.

2. Methods

2.1. Study design and data collection

Data were collected from a university dormitory building as part of the Center for Research in Indoor and Personal Technologies (CRIPT) study. The study monitored indoor CO₂ concentrations and occupancy patterns in student rooms over a 68-day period from February 4 to April 13, 2025.

2.1.1. CO₂ monitoring

Indoor CO₂ concentrations were measured at one-minute intervals using nondispersive infrared (NDIR) sensors deployed in each monitored room. A total of 78 sensors recorded approximately 1.4 million minute-level observations across 138 data files. Of these, 12 sensors were matched to rooms with concurrent occupancy ground truth, providing the dataset for method development and validation.

The 12 sensors monitored rooms of two types: 9 single-occupancy rooms and 3 shared rooms (two assigned residents). Sensors were placed at desk height to capture the breathing zone CO₂ concentration representative of the occupied microenvironment.

2.1.2. Occupancy ground truth

Self-reported occupancy data were collected from 24 participants (two per single room, two per shared room) through a structured survey instrument. Each single-occupancy room has one assigned resident; however, two study participants were enrolled per room to provide independent survey responses. The CO₂ sensor measures the room environment and cannot distinguish individual




Figure 1: Overview of the proposed method. CO₂ time series and block-level occupancy labels are processed through four stages: (1) per-sensor AR(1) parameter fitting, (2) empty-room calibration, (3) block-level excess-to-duration conversion, and (4) Monte Carlo uncertainty quantification. The output is an estimated number of occupied minutes per four-hour block with an 80% credible interval.

occupants, so the CO₂-based duration estimate for a given block is the same for all participants associated with a sensor. Participant-level totals differ only because they are aggregated over different subsets of blocks based on each participant's self-reported presence schedule. Participants reported their presence or absence for each four-hour block of the day (02:00–06:00, 06:00–10:00, 10:00–14:00, 14:00–18:00, 18:00–22:00, 22:00–02:00), yielding a binary label for each block: 1 (present for some or all of the block) or 0 (absent for the entire block).

The ground truth thus provides block-level *presence* labels but not the exact number of minutes present within each block. This distinction is important: a label of 1 indicates that the subject was present for at least some portion of the four-hour block, which could range from a few minutes to the full 240 minutes.

2.1.3. Dataset summary

After matching CO₂ sensors to subjects with occupancy labels, the analysis dataset comprises 791 four-hour blocks: 539 blocks labeled as present (label = 1) and 252 blocks labeled as absent (label = 0). Of these, 774 blocks (535 present, 239 absent) had overlapping CO₂ data and produced duration estimates; the remaining 17 blocks were excluded because the sensor data did not cover the block's time window. Table 1 summarizes the dataset characteristics.

Table 1: Dataset summary statistics.

Characteristic	Value
Monitoring period	Feb 4 – Apr 13, 2025
Duration (days)	68
Number of sensors	12
Single-occupancy rooms	9
Shared rooms	3
Number of subjects	24
Total four-hour blocks	791
Present (label = 1)	539
Absent (label = 0)	252
CO ₂ minute-level observations	$\sim 1.4 \times 10^6$

2.2. CO₂ dynamics model

We model indoor CO₂ concentration using a discrete-time autoregressive AR(1) process that captures the essential physics of indoor CO₂ generation and removal. Let C_t denote the smoothed CO₂ concentration at minute t and C_b denote the ambient baseline concentration. The excess CO₂ above baseline is defined as:

$$E_t = \max(C_t - C_b, 0) \tag{1}$$

The AR(1) dynamics relate consecutive excess values:

$$E_t = \phi E_{t-1} + g \cdot O_t + \varepsilon_t \tag{2}$$

where $\phi \in (0, 1)$ is the decay parameter (related to the air exchange rate), $g > 0$ is the CO₂ generation rate per minute of occupancy (ppm/min), $O_t \in \{0, 1\}$ is the occupancy state at minute t , and $\varepsilon_t \sim \mathcal{N}(0, \sigma_\varepsilon^2)$ is the innovation noise.

2.2.1. Physical interpretation of model parameters

The decay parameter ϕ captures the fraction of excess CO₂ that persists from one minute to the next. It is related to the effective air exchange rate λ (in air changes per minute) by $\phi = 1 - \lambda$. For typical dormitory rooms with natural or mechanical ventilation, λ ranges from 0.005 to 0.03 per minute (0.3 ACH to 1.8 ACH), corresponding to $\phi \in [0.97, 0.995]$.

The generation rate g represents the CO₂ emission rate of the room's occupant(s) expressed as the concentration increase per minute of presence. At steady state with continuous single-person occupancy, the excess CO₂ converges to:

$$E_{ss} = \frac{g}{1 - \phi} \quad (3)$$

For example, with $g = 15$ ppm/min and $\phi = 0.995$, the steady-state excess would be $15/0.005 = 3000$ ppm above baseline—though in practice, occupancy is intermittent and steady state is rarely achieved.

The innovation noise σ_ε captures sensor measurement error, short-term CO₂ fluctuations from sources other than occupancy (cooking, open windows, corridor air mixing), and model misspecification.

2.2.2. Baseline estimation

The ambient baseline C_b is estimated as the 1st percentile of all CO₂ readings for each sensor across its entire monitoring period. This captures the true outdoor/ambient air level (~ 400 ppm to 420 ppm) rather than the potentially elevated concentrations during nominally “empty” periods when residual CO₂ from prior occupancy may still be present. A floor of 380 ppm is enforced to ensure physical plausibility.

2.3. Parameter estimation

For each sensor, we estimate ϕ , g , and σ_ε from the minute-level CO₂ time series using the known-empty (label = 0) and known-occupied (label = 1) periods.

2.3.1. Room-level occupancy labels for shared rooms

For single-occupancy rooms, the subject-level occupancy label directly indicates whether the room is empty. For shared rooms (two subjects per sensor), the subject-level label reflects individual presence, not room-level emptiness: a minute where subject A is absent but subject B is present still has elevated CO₂. We therefore compute a room-level occupancy indicator for each sensor-minute pair:

$$O_t^{\text{room}} = \max_{j \in \mathcal{J}_s} O_{j,t} \quad (4)$$

where \mathcal{J}_s is the set of subjects mapped to sensor s . All physics fitting (decay rate, noise, empty-room floor) uses $O_t^{\text{room}} = 0$ to identify truly-empty minutes, preventing label contamination in shared rooms.

2.3.2. Decay parameter (ϕ)

The decay parameter is estimated from CO₂ decay episodes during known-empty periods. We select consecutive minute pairs where (i) the room is labeled absent, (ii) the time gap is exactly one minute (no data gaps), and (iii) the preceding excess exceeds 30 ppm (ensuring we fit genuine CO₂ decay rather than noise near ambient). The ordinary least squares (OLS) estimator is:

$$\hat{\phi}_{\text{OLS}} = \frac{\sum_{t \in \mathcal{D}} E_{t-1} E_t}{\sum_{t \in \mathcal{D}} E_{t-1}^2} \quad (5)$$

where \mathcal{D} denotes the set of valid decay observations. The standard error is:

$$\text{SE}(\hat{\phi}) = \sqrt{\frac{\hat{\sigma}_\varepsilon^2}{\sum_{t \in \mathcal{D}} E_{t-1}^2}}, \quad \hat{\sigma}_\varepsilon^2 = \frac{1}{|\mathcal{D}| - 1} \sum_{t \in \mathcal{D}} (E_t - \hat{\phi}_{\text{OLS}} E_{t-1})^2 \quad (6)$$

With one-minute sampling, OLS estimates of ϕ are biased upward toward 1.0 due to extreme autocorrelation. We apply Bayesian shrinkage toward a physical prior ($\phi_{\text{prior}} = 0.97$, corresponding to ~ 0.03 air changes per minute):

$$\hat{\phi} = \lambda_s \phi_{\text{prior}} + (1 - \lambda_s) \hat{\phi}_{\text{OLS}}, \quad \lambda_s = \frac{w_{\text{prior}}}{w_{\text{prior}} + |\mathcal{D}|} \quad (7)$$

where $w_{\text{prior}} = 180$ is the effective prior sample size. The final estimate is clipped to $[\phi_{\text{min}}, \phi_{\text{max}}] = [0.85, 0.995]$.

2.3.3. Generation rate (g)

The generation rate is estimated using the steady-state relationship between excess CO₂ and occupancy. During occupied periods, the 75th percentile of excess CO₂ serves as a proxy for sustained-presence conditions:

$$\hat{g} = \max \left(\left(E_{\text{occ}}^{(75)} - \tilde{E}_{\text{empty}} \right) \cdot (1 - \hat{\phi}), \hat{g}_{\text{innov}}, 1.0 \text{ ppm/min} \right) \quad (8)$$

where $E_{\text{occ}}^{(75)}$ is the 75th percentile of excess during occupied minutes, \tilde{E}_{empty} is the median excess during empty minutes, and \hat{g}_{innov} is an alternative estimate from the quantile of positive innovations during occupied periods. The generation rate is clipped to [1.0, 60.0 ppm/min].

2.3.4. Innovation noise (σ_ε)

The innovation noise is estimated as the standard deviation of innovations during known-empty, consecutive-minute periods:

$$\hat{\sigma}_\varepsilon = \text{SD} \left(\{E_t - \hat{\phi} E_{t-1} : t \in \mathcal{E}, \Delta t = 1\} \right) \quad (9)$$

where \mathcal{E} is the set of known-empty observations. This is clipped to [5, 200 ppm].

2.3.5. Fallback parameter rules

When a sensor has insufficient data for reliable parameter estimation, fallback values are used:

- If $|\mathcal{D}| < 30$ (fewer than 30 valid decay observations): $\hat{\phi} = \phi_{\text{prior}} = 0.97$, $\text{SE}(\hat{\phi}) = 0.03$.
- If $|\mathcal{O}| < 20$ (fewer than 20 occupied-period observations): \hat{g} defaults to the room-type mean (10 ppm/min for single rooms, 18 ppm/min for shared rooms), $\text{SE}(\hat{g}) = 4.0$.

Sensor 20233681 uses the fallback for $\hat{\phi}$ ($n_{\text{decay}} = 0$) and its generation rate $\hat{g} = 60$ ppm/min reflects the upper clipping bound rather than sensor-specific fitting. Table 2 identifies which sensors use fallback parameters.

2.4. Block-level excess estimation

The core estimation method operates at the four-hour block level, converting mean excess CO₂ within each block to estimated occupied minutes. The method consists of three stages: per-sensor empty-room calibration, self-normalizing scale computation, and occupied minutes estimation.

2.4.1. Per-sensor empty-room calibration

For each sensor, we compute block-level mean excess CO₂ for all blocks and establish an “empty-room floor” from the known-empty (label = 0) blocks. The empty-room floor captures residual CO₂ from prior occupancy, sensor offset, and building HVAC effects:

$$F_{\text{empty}} = \bar{E}_{\text{empty}} + k \cdot s_{\text{empty}} \quad (10)$$

where \bar{E}_{empty} and s_{empty} are the mean and standard deviation of block-mean excess across known-empty blocks, and k is a floor multiplier that controls the boundary between residual CO₂ and genuine occupancy signal.

Selection of floor multiplier (k).. The floor multiplier determines how aggressively residual CO₂ is suppressed. Because the finite-window normalization (Eq. 11) maps a given excess to more inferred minutes than steady-state normalization would, the empty-room floor must be raised correspondingly to maintain empty-room specificity. We selected k by a pre-specified grid search over $k \in [0.67, 1.00]$ in increments of 0.02, evaluating the calibration diagnostics C1 (< 15 min) and C2 ($> 60\%$). We chose the *smallest* k satisfying both criteria while preserving non-degenerate sensitivity, operationalized as a substantial fraction of label = 1 blocks retaining a positive estimated duration. This procedure selected $k = 0.86$; at this value, 47% of occupied blocks produce nonzero duration estimates.

2.4.2. Self-normalizing scale

Converting excess CO₂ to occupied fraction requires a normalization scale representing the expected block-mean excess during full continuous occupancy. Because occupancy blocks are finite (240 minutes) and steady state is rarely reached at high $\hat{\phi}$, we use a finite-window expected mean excess rather than the steady-state value $g/(1 - \phi)$. For an AR(1) process with continuous occupancy over n minutes, the expected block-mean excess is:

$$\bar{E}_{\text{full}}(n) = \frac{\hat{g}}{1 - \hat{\phi}} \left[1 - \frac{\hat{\phi}(1 - \hat{\phi}^n)}{n(1 - \hat{\phi})} \right] \quad (11)$$

We then blend this physics-based scale with a data-driven component:

$$S = \frac{1}{2} \left(E_{\text{occ}}^{(75)} - F_{\text{empty}} \right) + \frac{1}{2} \cdot \bar{E}_{\text{full}}(n) \quad (12)$$

where $E_{\text{occ}}^{(75)}$ is the 75th percentile of block-mean excess during known-occupied blocks and $n = 240$ is the block duration in minutes. The finite-window correction is substantial at high $\hat{\phi}$: for $\hat{\phi} = 0.995$ and $n = 240$, \bar{E}_{full} is approximately 42% of the steady-state value $g/(1 - \phi)$. The data-driven component makes cross-sensor estimates comparable even when \hat{g} and $\hat{\phi}$ differ substantially between sensors, while the physics component prevents degenerate cases.

2.4.3. Occupied minutes estimation

For each four-hour block, the calibrated excess is computed and converted to estimated occupied minutes:

$$E_{\text{cal}}^{(b)} = \max(\bar{E}^{(b)} - F_{\text{empty}}, 0) \quad (13)$$

$$f^{(b)} = \min\left(\frac{E_{\text{cal}}^{(b)}}{S}, 1\right) \quad (14)$$

$$\hat{M}^{(b)} = f^{(b)} \times 240 \quad (15)$$

where $\bar{E}^{(b)}$ is the mean excess for block b , $E_{\text{cal}}^{(b)}$ is the calibrated excess after subtracting the empty-room floor, $f^{(b)}$ is the estimated occupied fraction (clipped to $[0, 1]$), and $\hat{M}^{(b)}$ is the estimated occupied minutes (out of 240 minutes per block). The complete estimation pipeline is illustrated in Fig. 2, and Fig. 3 shows the model physics on example data.

2.5. Monte Carlo uncertainty quantification

Uncertainty in the duration estimates arises primarily from parameter estimation error (in $\hat{\phi}$, \hat{g} , and F_{empty}) and model misspecification. We propagate parameter uncertainty using Monte Carlo simulation with $N_{\text{MC}} = 200$ samples, treating the observed block means as fixed data (conditional inference).

For each Monte Carlo sample $s = 1, \dots, N_{\text{MC}}$:




Figure 2: Estimation pipeline for a single four-hour block. Raw CO₂ is converted to excess above baseline, averaged over the block, calibrated against the per-sensor empty-room floor, normalized by the blended scale S , and converted to occupied minutes.




Figure 3: Illustration of the CO₂ dynamics model and block-level estimation. The top panel shows the CO₂ time series with excess above baseline. The estimation procedure computes mean excess per block, subtracts the per-sensor empty-room floor, and normalizes by the blended scale S to obtain occupied fraction.

1. Draw perturbed parameters:

$$\phi^{(s)} \sim \text{TruncNormal}\left(\hat{\phi}, \text{SE}(\hat{\phi}), [\phi_{\min}, \phi_{\max}]\right) \quad (16)$$

$$g^{(s)} \sim \text{TruncNormal}(\hat{g}, \text{SE}(\hat{g}), [g_{\min}, g_{\max}]) \quad (17)$$

$$F^{(s)} \sim \text{TruncNormal}(F_{\text{empty}}, 0.5 s_{\text{empty}}, [0, \infty)) \quad (18)$$

2. Compute perturbed normalization scale $S^{(s)}$ using Eq. (12) with $\hat{g} \rightarrow g^{(s)}$ and $\hat{\phi} \rightarrow \phi^{(s)}$.
3. Compute estimated minutes $\hat{M}^{(b,s)}$ using Eqs. (13)–(15) with perturbed quantities and the observed (fixed) block-mean excess $\bar{E}^{(b)}$.

No separate observation-noise term is added to the block mean. The original formulation used $\eta \sim \mathcal{N}(0, \hat{\sigma}_\varepsilon^2/n_b)$, which assumes independent observations within a block. Because the AR(1) autocorrelation ($\hat{\phi} \approx 0.995$) severely violates this assumption, the effective sample size is far smaller than n_b and the σ/\sqrt{n} formula underestimates block-mean variance. Rather than substitute an ad-hoc empirical scale, we treat the observed block mean as fixed data (conditional inference) and let the parameter perturbation $(\phi, g, F_{\text{empty}})$ capture the dominant uncertainty sources.

The point estimate for each block is taken as the median across Monte Carlo samples ($\hat{M}_{\text{p}50}^{(b)}$), and the 80% credible interval is $[\hat{M}_{\text{p}10}^{(b)}, \hat{M}_{\text{p}90}^{(b)}]$. The uncertainty width is defined as:

$$W^{(b)} = \hat{M}_{\text{p}90}^{(b)} - \hat{M}_{\text{p}10}^{(b)} \quad (19)$$

Blocks are classified as “high confidence” when $W^{(b)} \leq 30 \text{ min}$ (i.e., the 80% credible interval spans no more than one eighth of the block duration).

2.6. Validation framework

We evaluate the method using two tiers of assessment criteria, illustrated in Fig. 4. The first tier comprises two *calibration diagnostics* (C1–C2) that were used to select the floor multiplier k (Section 2.4) and therefore cannot serve as independent validation. The second tier comprises four *independent evaluation criteria* (E1–E4) that test aspects of estimation quality not involved in any tuning decision. We additionally report two supplementary out-of-sample checks (C1-OOS, C2-OOS) with relaxed thresholds.

2.6.1. Calibration diagnostics (C1–C2)

These criteria were used to select the floor multiplier k (Section 2.4) and are therefore met by construction. They verify that the selected k achieves adequate empty-room specificity but do not constitute independent validation.

C1. Label = 0 mean unclamped minutes (<15 min): For blocks where the occupant self-reported absence for the entire four-hour period, the mean estimated duration should be small. “Unclamped” means estimates are computed using the same model without forcing label = 0 blocks to zero.




Figure 4: Validation framework organized into calibration diagnostics (C1–C2, green) and independent evaluation criteria (E1–E4, blue). C1–C2 were used to select the floor multiplier k and are therefore in-sample by construction. E1–E4 test cross-sensor generalizability, parameter robustness, and uncertainty informativeness independently of the calibration.

- C2. **Label = 0 % under 5 min (>60%)**: The majority of known-empty blocks should produce near-zero duration estimates. This is a stronger test than the mean criterion, requiring consistent (not just average) recognition of empty conditions.

2.6.2. Independent evaluation criteria (E1–E4)

These criteria were not involved in any tuning decision and provide independent assessment of estimation quality.

- E1. **Leave-one-sensor-out (LOO) mean relative MAE (<1.5)**: For each sensor, we estimate block durations using physics parameters ($\hat{\phi}, \hat{g}$) from the *remaining* sensors (cross-sensor median). The mean absolute error relative to sensor-specific estimates should be moderate, demonstrating that the method generalizes across rooms with different ventilation characteristics.
- E2. **LOO mean Spearman rank correlation (>0.2)**: The rank ordering of block durations should be preserved when using cross-sensor parameters. This tests whether the method captures *relative* occupancy patterns even when absolute calibration differs.
- E3. **Max ϕ sensitivity (% change) (<200%)**: Varying the decay parameter across the range [0.90, 0.99] should produce smooth, bounded changes in mean duration estimates. This tests robustness to the most uncertain physical parameter.
- E4. **% blocks with high confidence (>40%)**: A substantial fraction of occupied blocks should have tight uncertainty bounds (80% CI width ≤ 30 min). This tests whether the method provides actionable uncertainty information.

2.6.3. Out-of-sample temporal split

Calibration diagnostics C1 and C2 use the same empty blocks both to fit the empty-room floor F_{empty} and to evaluate it, creating an in-sample bias risk. To

verify that the calibration generalizes to unseen days, we repeat C1–C2 using a temporal split: for each sensor, empty blocks are sorted chronologically, the first 70% are used to fit F_{empty} , and the remaining 30% are evaluated with that floor. The out-of-sample criteria are relaxed to C1-OOS < 20 min and C2-OOS > 50%.

2.6.4. Baseline comparators

To demonstrate that the block-excess approach adds value beyond simpler heuristics, we compare against two baseline methods applied to the same sensor data:

- **Threshold method:** For each block, count the number of minutes where excess CO₂ exceeds a fixed threshold (50, 100, or 200 ppm).
- **Slope method:** For each block, count the number of minutes where CO₂ concentration is rising (positive first difference).

Each baseline is evaluated on the same C1 and C2 criteria as the block-excess estimator.

2.6.5. Semisynthetic validation

Because no ground-truth occupied minutes are available, we construct a semisynthetic validation by injecting known synthetic occupancy patterns into real empty-room CO₂ segments. The procedure is:

1. Select contiguous truly-empty blocks (≥ 200 minutes of data, mean excess < 20 ppm, room-level empty).
2. Define synthetic occupancy scenarios: continuous presence for 30, 60, 120, 180, or 240 minutes, fragmented patterns (4×30 min with gaps, 2×60 min with gap), and model-mismatch stress tests (see below).
3. Inject a *deterministic* synthetic excess into the real empty data using the sensor's fitted $\hat{\phi}$ and \hat{g} :

$$E_t^{\text{inj}} = \hat{\phi} E_{t-1}^{\text{inj}} + \hat{g} \cdot O_t^{\text{synth}} \quad (20)$$

The injected component contains no added noise because the real empty segment already carries realistic innovations and building transients; adding noise would double-count. The synthetic CO₂ is then $C_t^{\text{synth}} = C_t^{\text{empty}} + E_t^{\text{inj}}$.

4. For stress tests, injection uses parameters that differ from those assumed by the estimator, testing robustness to model mismatch: generation-rate mismatch ($\hat{g}' = 0.7\hat{g}$ and $1.3\hat{g}$), decay-rate drift ($\hat{\phi}' = \hat{\phi} + 0.005$), a +30 ppm baseline shift, and extra innovation noise ($\eta_t \sim \mathcal{N}(0, \hat{\sigma}_\varepsilon^2)$ added to the injected component).
5. Run the block-excess estimator on the synthetic data using the same F_{empty} and scale S as the real estimation.
6. Compute MAE, bias, and error distributions across scenarios.

3. Results

3.1. Sensor physics parameters

Table 2 presents the fitted AR(1) model parameters for each of the 12 sensors. The decay parameter $\hat{\phi}$ ranges from 0.970 to 0.995, corresponding to effective air exchange rates of 0.005 to 0.030 per minute (0.3 ACH to 1.8 ACH). The CO₂ generation rate \hat{g} ranges from 5.4 ppm/min to 60.0 ppm/min, reflecting differences in room size, ventilation configuration, and occupant metabolic characteristics. Baseline CO₂ concentrations range from 405.8 ppm to 542.0 ppm, with most sensors near the typical outdoor level of \sim 410 ppm and one sensor (22168412) showing an elevated baseline of 542.0 ppm, likely reflecting reduced ventilation or proximity to a CO₂ source. Sensor 20233681 has zero valid decay observations and uses the prior default for $\hat{\phi}$ (flagged in Table 2).

Table 2: Fitted AR(1) model parameters for each sensor. C_b : baseline CO₂ concentration; $\hat{\phi}$: decay parameter; \hat{g} : generation rate; $\hat{\sigma}_\varepsilon$: innovation noise standard deviation; n_{decay} : number of valid decay observations; Fallback: parameters using prior defaults instead of sensor-specific fits.

Sensor ID	Room type	C_b (ppm)	$\hat{\phi}$	\hat{g} (ppm/min)	$\hat{\sigma}_\varepsilon$ (ppm)	n_{decay}	Fallback
20233647	Single	443.7	0.992	18.9	17.1	693	none
20233665	Single	481.4	0.995	30.5	31.6	1658	none
20233681	Single	409.0	0.970	60.0	35.0	0	$\hat{\phi}$
20233685	Single	457.1	0.991	14.5	7.3	2081	none
20251591	Single	405.9	0.987	11.9	21.0	1897	none
22165976	Shared	405.8	0.995	16.3	9.0	2880	none
22165978	Single	477.3	0.994	5.4	9.7	2974	none
22165989	Single	419.3	0.995	22.8	15.1	1200	none
22165990	Single	412.6	0.970	26.2	5.7	50	none
22165993	Single	475.6	0.987	30.2	11.9	240	none
22165994	Shared	407.4	0.987	7.7	10.1	353	none
22168412	Shared	542.0	0.984	19.2	22.4	240	none

Fig. 5 visualizes the distribution of sensor parameters across rooms. The strong clustering of $\hat{\phi}$ near 0.995 reflects the tight, well-insulated construction typical of modern dormitory buildings.




Figure 5: Distribution of fitted sensor physics parameters across the 12 monitored rooms.

3.2. Block-level duration estimates

Table 3 summarizes the estimated occupancy duration by occupancy label. For known-absent blocks (label = 0), the unclamped estimates average 14.9 min with a block-level median of 0.0 min, matching the C1 calibration diagnostic exactly (both computed from the same unclamped re-estimation). Overall, 85.4% of absent blocks estimate under 5 min, confirming that the method correctly identifies empty rooms in the large majority of cases. For downstream use, label = 0 blocks are clamped to zero. For known-present blocks (label = 1), the mean estimated duration is 69.4 min with a median of 0.0 min; 47% of occupied blocks produce nonzero estimates. This right-skewed distribution reflects the conservative empty-room floor and the nature of four-hour block labeling: many blocks labeled “present” involve brief visits that do not elevate CO₂ above the per-sensor floor, while a smaller number of blocks show sustained presence with correspondingly higher duration estimates.

Table 3: Block-level duration estimates by occupancy label. Label = 0 values are *unclamped* (not forced to zero); for downstream use they are clamped. All values in minutes (out of 240-minute blocks). *n*: number of blocks; uncertainty refers to the width of the 80% credible interval.

Label	<i>n</i>	Mean	Median	SD	Min	Max	Mean uncertainty
0 (absent, unclamped)	239	14.9	0.0	47.7	0.0	240.0	28.9
1 (present)	535	69.4	0.0	90.1	0.0	240.0	61.2

The right-skewed distribution of label = 1 estimates (Fig. 6) is consistent with the nature of four-hour block labeling: a “present” label indicates the occupant was in the room for *at least some portion* of the block. Many blocks labeled as present may correspond to brief entries (returning to retrieve an item, checking in briefly) that produce modest CO₂ elevation above the empty-room floor, while extended study or sleep sessions produce sustained elevation and correspondingly higher duration estimates.

3.3. Time-of-day patterns

Fig. 7 shows the estimated occupancy duration by time-of-day block. As expected for a university dormitory population, occupancy is highest during the overnight block (22:00–02:00) when students are sleeping and lowest during the midday blocks (10:00–14:00 and 14:00–18:00) when students are attending classes and engaged in campus activities.

3.4. CO₂ trace examples

Fig. 8 presents representative CO₂ time series for selected sensors, illustrating the relationship between CO₂ dynamics and occupancy patterns. The traces demonstrate the characteristic rise during occupancy (CO₂ generation from breathing) and exponential decay after departure (ventilation removing excess CO₂). The variation in baseline levels, peak concentrations, and decay




Figure 6: Distribution of estimated occupancy duration by label. Label = 0 blocks cluster at zero minutes. Label = 1 blocks show a right-skewed distribution with many low-duration blocks and a tail extending to 240 minutes.




Figure 7: Estimated occupancy duration by time-of-day block. Error bars show the interquartile range across subjects. The overnight block (22:00–02:00) shows the highest occupancy, while midday blocks show the lowest, consistent with university student schedules.

rates across sensors reflects the heterogeneity of room ventilation and occupant behavior captured by the per-sensor parameter fitting.




Figure 8: Representative CO₂ time series for selected sensors. Shaded regions indicate four-hour blocks labeled as present (label = 1). The exponential decay after occupancy departure and the rise during presence illustrate the AR(1) dynamics captured by the model.

3.5. Validation results

Table 4 presents the validation results. Both calibration diagnostics (C1–C2) are met by construction (these criteria were used to select k). All four independent evaluation criteria (E1–E4) are satisfied. The out-of-sample temporal split test (C1-OOS) exceeds its criterion, indicating that the empty-room calibration benefits from the full dataset. These results are discussed in Section 4.

3.5.1. Label = 0 sanity check

The mean unclamped estimated duration for the 239 known-empty blocks is 14.9 min, below the 15 min criterion. Of these blocks, 85.4% estimate less than 5 min of occupancy (82.4% estimate exactly zero), with a heavy concentration near zero and a tail of larger residual-duration false positives for the remaining 14.6% of blocks (Fig. 9).

The 14.6% of empty blocks that estimate above 5 min represent cases where residual CO₂ from occupancy in a preceding block has not fully decayed during the current block, producing excess CO₂ that exceeds the per-sensor empty-room floor. This is physically expected given the long decay time constants ($\tau = 1/(1-\phi) \approx 33$ min to 200 min) and is captured by the uncertainty bounds.

Table 4: Validation results. Calibration diagnostics (C1–C2) are met by construction; all four independent evaluation criteria (E1–E4) are satisfied. Out-of-sample C1-OOS exceeds its criterion.

Tier	Test	Value	Criterion	Pass?	Interpretation
<i>Calibration diagnostics (used to select k; in-sample by construction)</i>					
C1:	Label = 0 mean unclamped min	14.9	<15	YES	Empty-room floor correctly calibrated
C2:	Label = 0 % under 5 min	85.4%	>60%	YES	Majority of absent blocks near zero
<i>Out-of-sample temporal split</i>					
	C1-OOS: mean unclamped min	46.0	<20	NO	Calibration degrades with limited data
	C2-OOS: % under 5 min	61.4%	>50%	YES	Majority of held-out absent blocks near zero
<i>Independent evaluation criteria</i>					
E1:	LOO mean relative MAE	1.11	<1.5	YES	Cross-sensor parameters within 150%
E2:	LOO mean Spearman ρ	0.313	>0.2	YES	Rank ordering preserved cross-sensor
E3:	Max ϕ sensitivity (%)	30.7%	<200%	YES	Smooth response to decay parameter
E4:	% blocks high confidence	50.7%	>40%	YES	About half of estimates have tight uncertainty




Figure 9: Distribution of unclamped estimated duration for known-empty (label = 0) blocks. The majority cluster at or near zero, with a tail of false positives from blocks where residual CO₂ from prior occupancy exceeds the per-sensor empty-room floor.

3.5.2. Leave-one-sensor-out cross-validation

Table 5 presents the leave-one-sensor-out results. For each sensor, block durations are estimated using the median $\hat{\phi}$ and median \hat{g} from the remaining 11 sensors (cross-sensor parameters). The mean absolute error relative to sensor-specific estimates is 1.11 (criterion: <1.5), and the mean Spearman rank correlation is 0.313 (criterion: >0.2).

Table 5: Leave-one-sensor-out cross-validation results. $\hat{\phi}_{\text{own}}$, \hat{g}_{own} : sensor-specific parameters; $\hat{\phi}_{\text{cross}}$, \hat{g}_{cross} : median of remaining sensors; MAE: mean absolute error (minutes); Rel. MAE: relative MAE (ratio); ρ : Spearman rank correlation between own and cross estimates.

Sensor ID	Room	$\hat{\phi}_{\text{own}}$	\hat{g}_{own}	\hat{g}_{cross}	MAE	Rel. MAE	ρ
20233647	Single	0.992	18.9	19.2	51.9	1.14	0.62
20233665	Single	0.995	30.5	18.9	54.9	2.52	0.56
20233681	Single	0.970	60.0	18.9	67.7	0.37	0.09
20233685	Single	0.991	14.5	19.2	80.1	0.61	0.07
20251591	Single	0.987	11.9	19.2	57.4	0.74	0.03
22165976	Shared	0.995	16.3	19.2	82.3	0.94	0.23
22165978	Single	0.994	5.4	19.2	8.7	1.27	0.76
22165989	Single	0.995	22.8	18.9	0.6	0.57	0.00
22165990	Single	0.970	26.2	18.9	79.4	0.77	0.05
22165993	Single	0.986	30.2	18.9	92.8	1.58	0.26
22165994	Shared	0.987	7.7	19.2	4.3	1.16	0.86
22168412	Shared	0.984	19.2	18.9	58.5	1.62	0.23
Mean				53.2	1.11	0.313	

Cross-sensor performance varies considerably across rooms. Several sensors retain moderate-to-strong rank agreement under cross-sensor parameters (e.g., 22165994 with $\rho = 0.86$ and 22165978 with $\rho = 0.76$), while others show weak agreement (e.g., 22165989 with $\rho = 0.00$ and 20251591 with $\rho = 0.03$). This spread reflects ventilation heterogeneity and the fact that many blocks lie near the empty-room floor, where small calibration differences determine whether a block estimates as nonzero. The heterogeneity motivates the per-sensor calibration approach.

3.5.3. Sensitivity analysis

Fig. 10 shows the sensitivity of estimated duration to the decay parameter ϕ across the range [0.90, 0.99]. The mean estimated duration varies from 55.4 min ($\phi = 0.99$) to 100.0 min ($\phi = 0.90$), compared to 80.0 min at the default $\phi = 0.97$. The maximum percentage change relative to the default is 30.7% (at $\phi = 0.99$)—well within the 200% criterion. The monotonic decrease with increasing ϕ is physically intuitive: higher ϕ (slower decay) produces larger normalization scale, so a given observed excess corresponds to a smaller occupied fraction.

Fig. 11 shows the sensitivity to the generation rate multiplier. Scaling the generation rate from 0.5× to 2.0× produces a monotonic decrease in estimated duration from 88.9 min to 52.2 min, consistent with the physics: higher gener-




Figure 10: Sensitivity of estimated occupancy duration to the decay parameter ϕ . Mean and median estimated duration (across all label = 1 blocks) are shown for each ϕ value. The smooth, monotonic response indicates robust behavior without discontinuities or instabilities.

ation rates produce more CO₂ per minute of occupancy, so a given observed excess corresponds to fewer occupied minutes.

3.6. Uncertainty analysis

Fig. 12 shows the distribution of uncertainty widths (80% credible interval spans) across all occupied blocks. The mean uncertainty width is 61.2 min, and 50.7% of occupied blocks achieve “high confidence” status (width \leq 30 min). High-confidence blocks tend to be those with either very low or very high mean excess CO₂: blocks near the empty-room floor produce tight estimates near zero, while blocks with very high excess produce tight estimates near the block maximum. Intermediate excess levels, where the calibrated signal is comparable to the parameter uncertainty, naturally produce wider confidence intervals.

3.7. Individual subject profiles

Table 6 reports participant-level summaries constructed from *room-level* CO₂-based duration estimates for the 18 participants in 9 single-occupancy rooms. Because each sensor measures a shared room environment, the CO₂-based duration estimate for a given block is the same for all participants associated with that sensor. Participant-level totals therefore differ only because they are aggregated over different subsets of blocks based on each participant’s self-reported presence (i.e., schedule-weighting), and should not be interpreted as CO₂-based disaggregation of individuals. For example, participants 324 and 325 (sensor 20233665) average 52 and 104 min/day respectively because they reported being present for different blocks. Participants with identical survey responses produce identical estimates (73/95, 506/507). The estimated




Figure 11: Sensitivity of estimated occupancy duration to the generation rate multiplier. Multipliers of $0.5\times$ to $2.0\times$ are applied to each sensor's fitted \hat{g} .




Figure 12: Uncertainty analysis of block-level duration estimates. Distribution of 80% credible interval widths across all occupied (label = 1) blocks.

mean daily room occupancy ranges from 0.0 min/day (participants 506, 507) to 985.5 min/day (participant 96).

Table 6: Participant-level occupancy summary for single-occupancy rooms (18 participants, 9 sensors). Each sensor monitors one room; two study participants per room provided independent self-reported presence schedules. The CO₂-based duration estimate is computed at the room level, so differences between participants on the same sensor arise solely from different self-reported presence patterns, not from distinct CO₂ signals. Mean, median, and standard deviation of daily estimated room occupancy (minutes per day). P10 and P90 are the mean daily 10th and 90th percentile estimates from the Monte Carlo uncertainty propagation.

Subject	Sensor	Days	Mean	Median	SD	P10	P90
96	20233681	7	985.5	1046.1	353.9	374.5	1199.3
151	20233681	7	937.8	814.1	230.7	346.7	1191.8
474	20233685	6	539.4	591.3	362.3	456.6	640.1
471	22165990	6	449.1	500.3	202.1	305.2	626.2
508	20233685	4	381.5	361.4	77.0	333.0	442.8
472	22165990	6	359.2	439.5	246.3	252.1	471.9
37	22165993	6	264.6	289.2	244.2	157.7	426.0
38	22165993	6	264.6	289.2	244.2	157.7	395.3
117	20233647	7	205.7	0.0	351.3	183.4	224.6
154	20233647	7	188.7	0.0	338.6	152.4	235.3
46	20251591	7	148.0	91.5	168.2	96.0	218.2
42	20251591	1	48.9	48.9	—	0.0	99.5
325	20233665	6	104.2	0.0	161.6	19.6	230.6
324	20233665	6	52.1	0.0	122.9	7.1	179.2
73	22165978	6	13.7	0.0	33.6	2.3	25.0
95	22165978	6	13.7	0.0	33.6	2.3	25.0
506	22165989	6	0.0	0.0	0.0	0.0	38.3
507	22165989	6	0.0	0.0	0.0	0.0	38.3

Table 7 presents supplementary room-level results for the 3 shared rooms (6 subjects). Because the CO₂ sensor cannot distinguish individual occupants, these estimates reflect *room-level* occupancy rather than per-person duration and are presented by sensor rather than by subject.

Table 7: Room-level occupancy for shared rooms (3 sensors, 6 subjects). Values reflect combined room occupancy and cannot be attributed to individual occupants.

Sensor	Subjects	Days	Mean	Median	P10	P90
22168412	132, 133	6	305.1	289.5	152.2	472.8
22165976	75, 76	6	173.0	85.5	71.4	308.9
22165994	399, 400	7	13.8	0.0	0.0	120.0

Fig. 13 visualizes the daily occupancy profiles for selected subjects, illustrating the diversity of temporal patterns captured by the method.




Figure 13: Individual subject daily occupancy profiles for selected subjects. Each bar represents estimated room occupancy for one day, with error bars showing the 80% credible interval. The diversity of patterns reflects different student lifestyles, from near-full-time room dwellers (top) to rarely present occupants (bottom).

3.8. Shared vs. single-occupancy rooms

Table 8 compares estimation results between shared and single-occupancy rooms. Single-occupancy rooms show higher mean estimated duration (78.6 min vs. 43.0 min) and higher mean CO₂ concentrations (1637 ppm vs. 1251 ppm). The lower duration estimates for shared rooms may reflect greater ventilation (more door openings with two occupants), measurement dilution (one sensor capturing a shared space), or genuinely lower per-person occupancy time in rooms with social dynamics.

Table 8: Comparison of block-level estimates between shared and single-occupancy rooms (label = 1 blocks only).

Room type	n_{blocks}	Subjects	Mean (min)	Median (min)	SD (min)	Mean CO ₂ (ppm)
Shared	138	6	43.0	0.0	71.9	1251
Single	397	18	78.6	5.8	94.1	1637

3.9. Daily heatmap

Fig. 14 presents a heatmap of estimated occupancy across subjects and days, providing a compact overview of the temporal patterns in the dataset. The heatmap reveals both consistent patterns (e.g., some subjects with regular nighttime occupancy) and sporadic patterns (e.g., subjects who are occasionally absent for extended periods), consistent with the variability of university student schedules.

4. Discussion

4.1. Out-of-sample empty-room calibration

To verify that the empty-room floor F_{empty} generalizes beyond the blocks used to fit it, we repeated the C1–C2 tests using a temporal split: the first 70% of each sensor’s empty blocks (chronologically) were used to fit F_{empty} , and the remaining 30% were evaluated with that floor. The out-of-sample results are reported alongside the in-sample results in Table 4. C2-OOS remains acceptable (61.4% of held-out empty blocks under 5 min), but C1-OOS increases substantially to 46.0 min, indicating that the empty-room floor is not fully stationary across days. This likely reflects day-to-day changes in residual CO₂ carryover (e.g., ventilation transients, atypical high-excess events, and changing occupancy patterns in adjacent blocks) that are not well captured when the calibration period is short. In practice, longer or periodically refreshed calibration windows may be needed for stable empty-room specificity.




Figure 14: Heatmap of estimated daily room occupancy (minutes) across subjects and days. Darker colors indicate higher estimated occupancy. Missing data (days without CO₂ records or occupancy labels) appear as gaps.

4.2. Comparison with baseline methods

Fig. 15 compares the block-excess estimator against simpler baselines: threshold methods (counting minutes where excess CO₂ exceeds 50, 100, or 200 ppm) and a slope method (counting minutes with rising CO₂). Because the floor multiplier k was selected using the empty-room calibration diagnostics (C1–C2), the block-excess method is expected to achieve strong empty-room specificity by design. The baseline comparators illustrate the trade-offs of fixed-threshold and slope heuristics: low thresholds inflate false positives in empty periods, high thresholds miss low-intensity occupancy, and the slope method is sensitive to noise. In contrast, the block-excess approach provides a per-sensor calibrated floor and yields interpretable duration estimates under a unified physics-based normalization.




Figure 15: Comparison of duration estimation methods on empty-room specificity. (a) C1: mean estimated minutes for known-empty blocks (lower is better). (b) C2: percentage of empty blocks estimating under 5 min (higher is better). (c) Mean estimated minutes for occupied blocks. The block-excess method achieves the calibration targets for empty-room specificity; baselines illustrate the trade-offs of simpler heuristics.

4.3. Semisynthetic validation

Because no ground-truth occupied minutes are available, we constructed a semisynthetic validation by injecting known synthetic occupancy patterns into real empty-room CO₂ segments (see Section 2.6.5). Fig. 16 shows the results and Table 9 provides numeric accuracy by scenario. Across continuous scenarios, the estimator shows systematic underestimation, with bias ranging from approximately -25 min (30–120 min cases) to -53 min for full-block (240 min) occupancy. Fragmented patterns produce comparable absolute errors but tend toward more negative bias than continuous occupancy of equal total duration (e.g., frag_4×30 bias = -48 min vs. cont_120 bias = -24 min), consistent with block-level averaging smoothing out short visits. The model-mismatch stress tests show that estimation degrades gracefully under parameter perturbation: generation-rate mismatch ($\hat{g} \times 0.7$ and $\times 1.3$), decay-rate drift ($\hat{\phi} + 0.005$), baseline shifts (+30 ppm), and extra innovation noise all produce bounded degradation (MAE 64 min to 115 min for a true duration of 120 min).




Figure 16: Semisynthetic validation results. (a) Scatter plot of estimated vs. true injected occupied minutes. Continuous scenarios (circles) cluster closer to the identity line than fragmented scenarios (triangles). (b) Distribution of absolute estimation error by true duration level.

Table 9: Semisynthetic validation accuracy by scenario. MAE: mean absolute error; Bias: mean signed error (positive = overestimation); P90: 90th percentile of absolute error. All values in minutes. Stress tests inject with parameters that differ from those used by the estimator.

Group	Scenario	True min	<i>n</i>	MAE	Bias	P90
Continuous	cont_30	30	11	24.5	-24.5	30.0
Continuous	cont_60	60	11	31.4	-27.6	60.0
Continuous	cont_120	120	11	63.3	-24.0	120.0
Continuous	cont_180	180	11	93.1	-33.5	180.0
Continuous	cont_240	240	9	52.7	-52.7	233.9
Fragmented	frag_4x30	120	11	49.5	-48.1	120.0
Fragmented	frag_2x60	120	11	50.9	-39.2	120.0
Stress	$g \times 0.7$	120	11	64.3	-64.3	120.0
Stress	$g \times 1.3$	120	11	103.7	+16.4	120.0
Stress	baseline	120	11	70.7	-16.6	120.0
Stress	$\phi + 0.005$	120	11	115.1	+33.4	120.0
Stress	extra noise	120	11	63.9	-22.6	120.0

This validation is “model-matched” in the sense that both injection and estimation use the same AR(1) model form. The stress tests partially address this limitation by perturbing generation rate, decay parameter, baseline level, and noise, but the fundamental model structure remains matched. A fully model-free external validation would require minute-level ground truth, which was not available in this study.

4.4. Strengths of the approach

The block-level excess method has several properties that make it well suited for CO₂-based duration estimation in residential settings.

Physics-based with minimal calibration. The method is grounded in well-understood CO₂ mass balance physics (Eq. 2) and requires only known-empty calibration periods—not labeled duration data—for per-sensor adaptation. This makes it deployable in settings where binary presence/absence labels are available (e.g., from access card data or simple surveys) but minute-level duration ground truth is not.

Block-level averaging exploits the SNR advantage. By working at the four-hour block level rather than the minute level, the method aggregates the cumulative CO₂ signal over a long window. While the AR(1) autocorrelation ($\hat{\phi} \approx 0.995$) limits the effective noise reduction below the i.i.d. \sqrt{n} rate, the block-level approach still transforms an intractable minute-level estimation problem into a tractable one. We verified this empirically by implementing a minute-level fused lasso deconvolution approach [18] that attempts to reconstruct the occupancy signal at one-minute resolution. The fused lasso estimator failed three of six validation criteria due to the fundamental SNR limitation: at $\phi \approx 0.995$, the minute-level occupancy innovation (5 ppm to 60 ppm) is comparable to the noise standard deviation (5 ppm to 40 ppm), yielding SNR of order unity. Block-level aggregation overcomes this by integrating the occupancy-driven excess over 240 minutes.

Per-sensor calibration handles heterogeneity. The per-sensor empty-room floor (Eq. 10) and self-normalizing scale (Eq. 12) adapt to each room’s ventilation, sensor placement, and CO₂ baseline, avoiding the need for uniform room physics assumptions.

Full uncertainty quantification. The 200-sample Monte Carlo propagation produces 80% credible intervals that honestly reflect the uncertainty from parameter estimation ($\hat{\phi}, \hat{g}$) and empty-room floor estimation (F_{empty}). The observed block mean is treated as fixed data, avoiding the need for an observation-noise term whose scale is ill-defined under strong AR(1) autocorrelation. The 50.7% high-confidence rate provides users with clear guidance on which estimates are most reliable.

4.5. Limitations

Block-level temporal resolution. The method estimates total occupied minutes within four-hour blocks but does not resolve the temporal pattern of

occupancy within each block. A block with 120 estimated minutes could correspond to two continuous hours or twelve 10-minute visits. Finer temporal resolution would require either higher SNR conditions (lower ϕ , higher g) or complementary sensor modalities.

Self-reported ground truth limitations. The occupancy labels are binary (present/absent for the whole block) and self-reported, introducing two limitations. First, we cannot validate the *absolute accuracy* of duration estimates for label = 1 blocks since the true number of minutes is unknown. Validation instead relies on internal consistency (label = 0 sanity, LOO, sensitivity). Second, self-report is subject to recall bias, potentially mislabeling some blocks.

Single-zone assumption. The AR(1) model treats each room as a well-mixed single zone. In practice, CO₂ gradients exist within rooms, and sensor placement affects the measured concentration. This is partially captured by the per-sensor calibration but not explicitly modeled.

Multiple-occupant ambiguity in shared rooms. In shared rooms, the CO₂ signal reflects the combined presence of both occupants. Although we compute room-level occupancy labels (Eq. 4) to prevent label contamination during physics fitting, the method estimates total room-level occupied person-minutes and cannot distinguish individual contributions. The 3 shared-room sensors (6 subjects) should be interpreted as supplementary to the primary analysis on 9 single-occupancy sensors (18 subjects).

Right-skewed label = 1 distribution. The median estimated duration for label = 1 blocks is 0.0 min, indicating that over half of “present” blocks produce estimates indistinguishable from zero. This reflects the combination of a conservative empty-room floor (which correctly suppresses false positives) and the nature of four-hour block labeling: many blocks labeled “present” involve brief visits that barely elevate CO₂ above the floor. This is plausible (students often enter briefly to retrieve items) but cannot be verified without minute-level ground truth.

4.6. Comparison to literature

Most CO₂-based occupancy studies focus on detection (binary or count) rather than duration. Candanedo and Feldheim [6] achieved 97% accuracy for binary detection using CO₂, temperature, and humidity with random forests, but did not estimate duration. Chen et al. [8] reviewed CO₂-based occupancy estimation broadly, noting that most methods target presence or count rather than time spent. Jin et al. [7] developed a virtual CO₂ sensor approach for commercial buildings but targeted occupancy count rather than duration per person. More recent deep learning methods [11, 12] achieve high detection accuracy but remain focused on instantaneous state rather than cumulative duration within time windows.

Bayesian and state-space approaches have brought principled uncertainty quantification to CO₂-based occupancy inference. Ebadat et al. [13] formulated occupancy estimation as a regularized deconvolution problem, recovering occupancy profiles from CO₂ using total variation regularization. Wang et al. [14]

combined Markov-based recurrent neural networks with Wi-Fi probe data for occupancy prediction, demonstrating the value of sequential models for temporal occupancy patterns. Both methods estimate instantaneous occupancy state rather than block-level duration, and both require either known ventilation rates or concurrent training data with ground-truth occupancy counts.

The closest methodological precedent is the CO₂ mass-balance inversion approach used in ventilation rate estimation [19, 20, 21]. These methods solve the mass balance equation for airflow rate given known occupancy. Our method inverts this relationship, solving for occupancy duration given estimated ventilation parameters. The key difference is that ventilation estimation assumes known, stable occupancy, while our problem has uncertain, time-varying occupancy. Szczerba et al. [22] similarly used CO₂ mass balance to estimate the number of occupants in office rooms, but their approach targets instantaneous count and requires continuous ventilation rate measurements.

The block-level temporal resolution of our method contrasts with the minute-level approaches above but aligns with the time-use survey literature, where occupancy is commonly reported in multi-hour blocks [23, 24]. Our approach bridges these scales: it uses minute-level CO₂ data as input but produces outputs at the temporal resolution of self-reported surveys, enabling direct comparison with questionnaire-based occupancy data.

Our validation framework—distinguishing calibration diagnostics from independent evaluation criteria—is, to our knowledge, the first structured approach to validating CO₂-based duration estimation. We hope it provides a useful benchmark for future methods.

5. Conclusions

We have presented a physics-based method for inferring occupancy duration from indoor CO₂ concentrations in university dormitory rooms. The method operates at the four-hour block level, aggregating the cumulative CO₂ signal over long windows to overcome the low signal-to-noise ratio inherent in minute-level CO₂ data. Per-sensor empty-room calibration and a self-normalizing scale handle the heterogeneity of room ventilation and sensor characteristics without requiring uniform physics assumptions.

Applied to 12 sensors monitoring 24 participants over 68 days, all four independent evaluation criteria are satisfied (cross-sensor generalizability E1–E2, parameter sensitivity E3, and uncertainty informativeness E4). The two calibration diagnostics used to select the empty-room floor multiplier are met by construction. An out-of-sample temporal split shows that the empty-room calibration partially generalizes (C2-OOS passes) but degrades when training data are limited (C1-OOS exceeds its criterion), suggesting that longer calibration periods improve generalization. A semisynthetic validation—injecting known occupancy patterns with model-mismatch stress tests into real empty-room data—provides partial external validation of duration recovery. Comparison against threshold and slope baselines illustrates the trade-offs of simpler heuris-

tics and confirms that the block-excess approach achieves the calibration targets for empty-room specificity.

The method is computationally efficient (under one minute on standard hardware), requires no supervised training data—only binary block-level presence labels for calibration—and provides full Monte Carlo uncertainty quantification. It is applicable to any residential or office setting where CO₂ is measured continuously and binary presence/absence labels are available for calibration.

Future work should address minute-level temporal resolution (potentially through complementary sensor fusion), multi-occupant disaggregation in shared rooms, and validation against minute-level ground truth in controlled experiments.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

The authors thank the CRIPT research team for data collection and the dormitory residents who participated in the study.

Data availability

The data and code supporting this study are available from the corresponding author upon reasonable request.

References

- [1] Frontczak, M., Wargocki, P., 2012. Literature survey on how different factors affect human comfort in indoor environments. *Building and Environment* 46 (4), 922–937.
- [2] Daisey, J.M., Angell, W.J., Apte, M.G., 2003. Indoor air quality, ventilation and health symptoms in schools: an analysis of existing information. *Indoor Air* 13 (1), 53–64.
- [3] Fisk, W.J., 2017. The ventilation problem in schools: literature review. *Indoor Air* 27 (6), 1039–1051.
- [4] Yan, D., O'Brien, W., Hong, T., Feng, X., Burak Gunay, H., Tahmasebi, F., Mahdavi, A., 2015. Occupant behavior modeling for building performance simulation: current state and future challenges. *Energy and Buildings* 107, 264–278.

- [5] Dong, B., Yan, D., Li, Z., Jin, Y., Feng, X., Fontenot, H., 2018. Modeling occupancy and behavior for better building design and operation—a critical review. *Building Simulation* 11 (5), 899–921.
- [6] Candanedo, L.M., Feldheim, V., 2016. Accurate occupancy detection of an office room from light, temperature, humidity and CO₂ measurements using statistical learning models. *Energy and Buildings* 112, 28–39.
- [7] Jin, M., Bekiaris-Liberis, N., Weekly, K., Spanos, C.J., Bayen, A.M., 2018. Occupancy detection via environmental sensing. *IEEE Transactions on Automation Science and Engineering* 15 (2), 443–455.
- [8] Chen, Z., Jiang, C., Xie, L., 2018. Building occupancy estimation and detection: a review. *Energy and Buildings* 169, 260–270.
- [9] Yang, J., Santamouris, M., Lee, S.E., 2016. Review of occupancy sensing systems and occupancy modeling methodologies for the application in institutional buildings. *Energy and Buildings* 121, 344–349.
- [10] Masood, M.K., Soh, Y.C., Jiang, C., 2019. Occupancy estimation from environmental parameters using wrapper and hybrid feature selection. *Applied Soft Computing* 60, 482–494.
- [11] Chen, Z., Zhao, R., Zhu, Q., Masood, M.K., Soh, Y.C., Mao, K., 2020. Building occupancy estimation with environmental sensors via CDBLSTM. *IEEE Transactions on Industrial Electronics* 67 (11), 9659–9668.
- [12] Tekler, Z.D., Low, R., Gunay, B., Andersen, R.K., Blessing, L., 2020. A scalable Bluetooth Low Energy approach to identify occupancy patterns and profiles in office spaces. *Building and Environment* 171, 106681.
- [13] Ebadat, A., Bottegal, G., Varagnolo, D., Wahlberg, B., Johansson, K.H., 2015. Estimation of building occupancy levels through environmental signals deconvolution. In: *Proceedings of the 5th ACM Workshop on Embedded Systems for Energy-Efficient Buildings*, pp. 1–8.
- [14] Wang, W., Chen, J., Hong, T., Zhu, N., 2018. Occupancy prediction through Markov-based feedback recurrent neural network (M-FRNN) algorithm with Wi-Fi probe technology. *Building and Environment* 138, 160–170.
- [15] Persily, A., de Jonge, L., 2017. Carbon dioxide generation rates for building occupants. *Indoor Air* 27 (5), 868–879.
- [16] ASHRAE, 2022. ANSI/ASHRAE Standard 62.1-2022: Ventilation and Acceptable Indoor Air Quality. American Society of Heating, Refrigerating and Air-Conditioning Engineers, Atlanta, GA.
- [17] Lu, T., Lü, X., Viljanen, M., 2011. A novel and dynamic demand-controlled ventilation strategy for CO₂ control and energy saving in buildings. *Energy and Buildings* 43 (9), 2499–2507.

- [18] Tibshirani, R., Saunders, M., Rosset, S., Zhu, J., Knight, K., 2005. Sparsity and smoothness via the fused lasso. *Journal of the Royal Statistical Society: Series B* 67 (1), 91–108.
- [19] Persily, A.K., 2006. Indoor carbon dioxide concentrations in commercial and institutional buildings. *Indoor Air* 7 (3), 158–169.
- [20] Li, Z., Dong, B., 2021. A new modeling approach for short-term prediction of occupancy in residential buildings. *Building and Environment* 202, 108066.
- [21] Calì, D., Matthes, P., Huber, K., Osterhage, T., Müller, D., 2015. CO₂ based occupancy detection algorithm: experimental analysis and validation for office and residential buildings. *Building and Environment* 86, 39–49.
- [22] Szczurek, A., Maciejewska, M., Teuerle, M., Wyłomańska, A., 2017. Method to characterize collective impact of factors on indoor air. *Physica A: Statistical Mechanics and its Applications* 468, 851–862.
- [23] Wilke, U., Haldi, F., Scartezzini, J.-L., Robinson, D., 2013. A bottom-up stochastic model to predict building occupants' time-dependent activities. *Building and Environment* 60, 254–264.
- [24] Richardson, I., Thomson, M., Infield, D., 2010. A high-resolution domestic building occupancy model for energy demand simulations. *Energy and Buildings* 40 (8), 1560–1566.